

FinIF: Can Your Model Follow Instructions in Financial Tasks?

Anonymous ACL submission

Abstract

Large Language Models (LLMs) are deployed in finance, where tasks are embedded in multi-stage workflows. Unlike open-domain question answering, financial workflows impose strict delivery specifications; the viability of a generated artifact depends not only on semantic quality but also on satisfying instruction-following (IF) constraints. We observe that even frontier models struggle with these constraints, leaving some outputs insufficient as executable and auditable downstream work products. To bridge this gap, we introduce FinIF, a framework for evaluating and improving financial instruction following through a scalable automated data construction pipeline. Aligned with the O*NET occupational database, FinIF structures the financial task lifecycle into five continuous workflows and extends real-world materials into FinIF-Test, a 300-item benchmark with 3,134 constraints, and FinIF-Train, a 1,064-example instruction-tuning dataset. FinIF-Test pairs deterministic Python rule checkers for 1,346 structural constraints with calibrated LLM judges for 1,788 semantic constraints. Empirical evaluations reveal a behavioral dichotomy: GPT-5.5 achieves 93.30% micro-average constraint satisfaction, yet only 58.67% strict item-level pass due to failures in citation discipline, quantitative verification, and operational thresholds. FinIF exposes a critical transition failure of LLMs from advanced analysts to reliable workflow executors, and provides scalable data resources for financial alignment.

1 Introduction

Large language models (LLMs) are increasingly used for financial analysis, compliance review, investment due diligence, reporting, and operations support. In these settings, a user usually provides source materials, asks for a specific work product, and states delivery requirements that determine whether the output can be reviewed or used downstream. For example, a model may need to prepare

a credit memo, suitability review, reconciliation exception report, or compliance escalation note while citing active source labels, separating facts from assumptions, showing calculation inputs, preserving thresholds, and avoiding unsupported approval or legal conclusions.

This setting exposes a gap in current evaluation. A model can be financially plausible yet operationally unusable if it ignores delivery constraints. A response that reaches the right high-level recommendation may still fail because citations appear only at the end, a calculation lacks source inputs, a risk label is softened, or an approval boundary is omitted. In high-stakes workflows, such failures are not cosmetic. They break audit trails, weaken reviewability, and can create compliance risk.

Existing benchmarks only partially cover this problem. General instruction-following benchmarks evaluate whether models obey verifiable or fine-grained constraints, but their tasks are usually not grounded in financial work products (Zhou et al., 2023; Jiang et al., 2024; Qin et al., 2024). Financial benchmarks evaluate financial knowledge, numerical reasoning, question answering, or business-domain task quality (Islam et al., 2023; Shah et al., 2022; Guo et al., 2025; Nie et al., 2025; Xie et al., 2024; Lu et al., 2025). However, they do not isolate whether the model follows explicit finance-specific delivery instructions. FinIF targets this missing axis: whether an LLM can produce a usable financial deliverable under stated workflow constraints.

Figure 1 illustrates the distinction. In one FinIF-Test item, GPT-5.5 produced an accurate and well-grounded regulatory-dependency benchmark note and received a quality score of 10/10. Nevertheless, it failed the item because the response exceeded the specified word range. This is a small example of the central issue: high output quality and instruction compliance are related but not interchangeable, especially when the output is expected to be audit-

Example FinIF-Test failure. A strong model writes a high-quality financial note, cites the supplied materials, and gives a useful next action. The evaluator still marks the item as not strictly passed because one explicit delivery requirement is violated: the response contains 539 words where the prompt requires 180–430 words.

Figure 1: A high-quality answer can still fail as a financial work product when an explicit delivery constraint is violated.

ready.

We present **FinIF**, a finance instruction-following benchmark suite. FinIF has two components. **FinIF-Test** is a 300-item hard benchmark selected from workflow-grounded finance tasks. It contains 3,134 instruction-following constraints, all evaluated with 100% coverage in our experiments. **FinIF-Train** contains 1,064 workflow-grounded examples with 7,851 annotated constraints, intended for supervised training, preference-data construction, and further benchmark expansion.

FinIF differs from prior work in three ways. First, each item is built from a concrete financial workflow node and work product, rather than from a generic prompt template. Second, the benchmark separates instruction-following constraints from auxiliary diagnostic axes such as answer quality and finance validity. Third, the evaluator combines rule-based checking for mechanical constraints with LLM-as-a-judge scoring for semantic constraints, while reporting strict item-level compliance in addition to micro constraint pass rate.

Our experiments reveal that current models remain far from fully reliable on financial instruction following. GPT-5.5 passes 93.30% of individual constraints but strictly passes only 58.67% of items. GPT-5 passes 91.35% of individual constraints but only 47.33% of complete items. Smaller or cheaper models show a larger gap, with DS-V4-Pro, DS-V4-Flash, and Qwen3.5-9B reaching strict item pass rates of 16.00%, 9.00%, and 4.00%, respectively. These results suggest that “mostly compliant” outputs are common, but complete compliant work products remain difficult.

Our main contributions are:

- We introduce FinIF-Test, a 300-item hard benchmark for evaluating instruction following in financial workflow tasks.
- We construct FinIF-Train, a 1,064-example workflow-grounded training pool with 7,851

annotated constraints.

- We propose a scalable construction pipeline and a hybrid evaluator that combines deterministic rule checks with item-batched LLM judging.
- We show that strong models still fail strict financial delivery requirements, especially evidence placement, quantitative verification, and decision-boundary reasoning.

2 Related Work

2.1 Instruction Following of LLMs

Instruction-following evaluation has moved from broad human preference scoring toward explicit requirement checking. IFEval introduced a reproducible benchmark of verifiable constraints such as length, keyword, and format requirements (Zhou et al., 2023). FollowBench extends this direction with multi-level fine-grained constraints (Jiang et al., 2024), while InFoBench decomposes instructions into atomic requirements for more interpretable scoring (Qin et al., 2024). LIFBench studies whether models can follow instructions reliably in long-context settings (Wu et al., 2025).

Recent work further shows that realistic settings require domain or scenario-specific constraint design. AgentIF evaluates long agentic instructions with tool, conditional, and example-implied constraints (Qi et al., 2025). SciIF argues that scientific answer correctness and scientific constraint compliance are distinct capabilities, and uses evidence-based judging to audit whether required constraints are explicitly satisfied (Su et al., 2026). The closest financial work is FIFE, which evaluates complex financial instruction following using 88 human-authored prompts and chainable verifiers (Matlin et al., 2025). FinIF follows this line but targets a different scale and use case: workflow-grounded financial deliverables with both a 300-item hard test set and a 1,064-example training pool.

2.2 Financial Capability Evaluation

Financial NLP benchmarks usually measure language understanding, knowledge, reasoning, question answering, or task quality rather than explicit instruction adherence. FLUE provides a financial language understanding suite across five NLP tasks (Shah et al., 2022). FinanceBench evaluates open-book financial question answering over public company materials with evidence strings (Islam

et al., 2023). FinEval and CFinBench evaluate Chinese financial knowledge and practical financial ability at large scale (Guo et al., 2025; Nie et al., 2025). FinBen aggregates financial tasks across information extraction, textual analysis, QA, generation, risk management, forecasting, and decision-making (Xie et al., 2024). BizFinBench emphasizes business-driven real-world financial applications and combines objective and subjective metrics (Lu et al., 2025).

These benchmarks are valuable for measuring whether models know finance or can solve finance-oriented tasks. However, an operational financial output must also satisfy delivery constraints: source placement, calculation traceability, threshold preservation, missing-evidence handling, and decision-boundary language. FinIF therefore treats instruction adherence as a first-class capability rather than a side effect of answer quality.

2.3 Automated and Auditable Evaluation

FinIF uses LLM judges only where deterministic rules are insufficient. This choice is motivated by prior work showing that LLM-as-a-judge can scale evaluation of open-ended outputs, while also requiring careful rubric design because judges can exhibit bias and inconsistent thresholds (Liu et al., 2023; Zheng et al., 2023). Accordingly, FinIF separates deterministic checks from semantic judging and reports strict item-level pass rate in addition to aggregate constraint accuracy. This design follows the auditability motivation in AgentIF and SciIF: a model should not receive credit for implicit compliance when the required evidence or boundary statement is absent from the output.

3 Proposed Benchmark: FinIF

FinIF evaluates whether an LLM can turn supplied financial materials into an instruction-compliant work product. Each item consists of a context pack, a workflow-specific query, a set of instruction-following constraints, and an evaluation protocol. The response is assessed along two separate axes: whether it follows the requested constraints, and whether it is a useful financial answer. The first axis is the main target of FinIF; the second is kept as an auxiliary diagnostic.

3.1 Task Taxonomy

FinIF uses five workflow stages that cover a broad finance task lifecycle: Intake and Profiling; Re-

Benchmark	Domain	Test	Train	Main evaluation
IFEval (Zhou et al., 2023)	General	541	–	Rule-based IF checks
AgentIF (Qi et al., 2025)	Agentic	707	–	Code, LLM, hybrid checks
SciIF (Su et al., 2026)	Science	334	910	Correctness + constraint compliance
FIIF (Matlin et al., 2025)	Finance	88	–	Chainable financial verifiers
FinIF	Finance	300	1,064	Rule + LLM audit for workflows

Table 1: Comparison with representative instruction-following benchmarks. FinIF focuses on financial workflow deliverables and provides both a hard test set and training data.

Workflow	Train	Test
Intake and Profiling	168	54
Research and Due Diligence	196	45
Decision and Structuring	224	63
Risk and Compliance Review	224	64
Execution, Monitoring, Reporting, and Operations	252	74
Total	1,064	300

Table 2: Workflow coverage of FinIF-Train and FinIF-Test.

search and Due Diligence; Decision and Structuring; Risk and Compliance Review; and Execution, Monitoring, Reporting, and Operations. These workflows are decomposed into 38 task types. Examples include KYC onboarding, loan application intake, financial statement analysis, credit due diligence, investment memo drafting, DCF valuation, suitability review, AML red-flag review, reconciliation, board or regulatory reporting, and remediation tracking. This workflow-first design keeps each prompt tied to a plausible financial deliverable rather than an isolated instruction template.

3.2 Data Construction

The FinIF construction pipeline follows a workflow-first sequence:

workflow → task → work product → context →
query → constraints → evaluation.

Context acquisition. We construct context packs from public primary sources, synthetic private records, or mixed packs. Public sources include regulatory guidance, public filings, financial statements, and similar primary materials. Private customer, borrower, account, and internal-policy records are synthetic to avoid exposing sensitive data while preserving realistic workflow triggers.

Query synthesis. Each item starts from a target work product, such as a KYC review note, credit memo, investment committee memo, suitability review note, compliance escalation report, reconciliation exception report, or board report. The query asks the model to produce this deliverable under explicit requirements about format, evidence,

quantitative reasoning, decision boundaries, and missing-information handling.

Constraint annotation. FinIF evaluates explicit, checkable requirements rather than implicit best practices. Each constraint receives a fine-grained tag, a family label, an evaluation axis, and a checker route. The paper-facing taxonomy contains five families: Format and Presentation, Evidence and Grounding, Decision and Boundary, Quantitative Verification, and Required Content Coverage. For evaluation, content-coverage constraints that primarily diagnose answer completeness are kept separate from instruction-following constraints when computing the main IF metrics.

Evaluation generation. For each constraint, we generate an evaluation protocol that routes mechanically checkable requirements to rule checkers and semantic requirements to an LLM judge. Rule-checked constraints include word ranges, decimal places, table columns, required labels, first-line formats, and ordered-list counts. LLM-judged constraints include source-label placement, fact-inference separation, evidence-rule-action chains, and boundary reasoning.

Quality control. Before inclusion, each item is validated for schema consistency, train/test separation, context availability, checker coverage, and parseable evaluator output. The reported FinIF-Test runs reach 100% decision coverage, meaning every instruction-following constraint is assigned a pass/fail decision.

3.3 Data Statistics

FinIF-Train contains 1,064 workflow-grounded examples across all 38 task types, with 28 examples per task and 7,851 annotated constraints. FinIF-Test contains 300 hard items, 3,134 instruction-following constraints, and 1,904 additional diagnostic constraints used for coverage, quality, and finance-validity analysis. The test set has an average of 10.45 instruction-following constraints per item, with a minimum of 7 and a maximum of 16. Among the 3,134 test IF constraints, 1,346 are rule-checked and 1,788 are LLM-judged.

4 Capability Evaluation

4.1 Experimental Setup

We evaluate five completed model runs on the same 300-item FinIF-Test split: GPT-5.5, GPT-5, DS-V4-Pro, DS-V4-Flash, and Qwen3.5-9B. Each run is

Constraint family	Test IF	Share
Format and Presentation	1,391	44.38%
Evidence and Grounding	1,085	34.62%
Decision and Boundary	280	8.93%
Quantitative Verification	377	12.03%
Required Content Coverage	1	0.03%
Total	3,134	100.00%

Table 3: Instruction-following constraint distribution in FinIF-Test. Content coverage is mainly used as an auxiliary diagnostic rather than the primary IF target.

Model	Micro IF	Macro IF	Strict	Quality
GPT-5.5	93.30	93.10	58.67	7.40
GPT-5	91.35	91.20	47.33	7.05
DS-V4-Pro	81.24	80.72	16.00	4.68
DS-V4-Flash	77.92	77.25	9.00	4.27
Qwen3.5-9B	72.08	71.45	4.00	3.60

Table 4: Main results on FinIF-Test. IF metrics are percentages. Quality is a 0–10 auxiliary judge score. Strict pass requires satisfying every IF constraint in an item.

evaluated against 3,134 instruction-following constraints. The evaluator reaches 100% decision coverage for all reported runs.

FinIF reports three primary instruction-following metrics. **Micro IF** is the fraction of passed constraints among all decided constraints. **Macro item IF** averages item-level constraint pass rates, giving each test item equal weight. **Strict item pass rate** counts an item as passed only if every instruction-following constraint in that item is satisfied. We also report a holistic 0–10 quality score from the LLM judge as an auxiliary diagnostic; it is not used as the primary instruction-following metric.

The evaluator follows a rule-first protocol. Mechanically checkable constraints are routed to deterministic checkers, including word ranges, decimal-place requirements, table columns, required labels, first-line formats, and ordered-list counts. Semantic constraints are batched by item and judged by GPT-4o with fixed prompts and strict JSON parsing. This hybrid design reduces unnecessary judge calls while preserving coverage for constraints that require semantic interpretation.

4.2 Main Results

Table 4 shows a consistent gap between local constraint compliance and complete work-product compliance. GPT-5.5 passes 2,924 of 3,134 IF

constraints, but strictly passes only 176 of 300 items. GPT-5 passes 2,863 constraints but strictly passes only 142 items. The gap is larger for cheaper and smaller models: DS-V4-Pro strictly passes 48 items, DS-V4-Flash 27 items, and Qwen3.5-9B only 12 items.

This gap is the central empirical finding of FinIF. Micro IF captures whether a model follows many local requirements, while strict item pass rate captures whether the response can be accepted as a complete financial work product. In production-like workflows, a response that violates one critical evidence, calculation, or boundary instruction may still require human repair even if most local constraints pass.

4.3 Further Analysis

Evidence placement is the dominant failure mode. Across the evaluated models, the most frequent failed tag is EG2, which requires active source labels next to material facts, figures, thresholds, missing-evidence triggers, and decision points. GPT-5 fails 119 EG2 constraints, and GPT-5.5 fails 107. This indicates that even strong models often cite sources globally but fail to place evidence exactly where an audit trail requires it.

Quantitative verification remains fragile. The second recurring weakness is QV2, which requires source inputs, formula or comparison, result, and business implication for calculations and reconciliations. GPT-5 fails 42 QV2 constraints, while DS-V4-Pro and Qwen3.5-9B fail 113 and 162, respectively. These failures are not necessarily arithmetic mistakes; many are lineage failures where the response gives a number without the required verification trail.

Decision-boundary reasoning is hard to keep connected. DB9 asks the response to connect active evidence, governing rule or task standard, and final action in a single reasoning chain. GPT-5.5 and GPT-5 still fail 27 and 22 DB9 constraints, respectively. Smaller models fail this tag much more often. This suggests that models can list facts and recommendations, but often struggle to maintain an auditable evidence-rule-action chain.

Format constraints are not solved. Although many format checks are mechanically simple, they continue to cause item failures. GPT-5.5 fails 19 FP3 constraints, and GPT-5 fails 58. These include length, decimal-place, count, and ordering

requirements. The Figure 1 example shows why this matters: an otherwise high-quality response can be rejected because it fails a concrete delivery constraint.

5 FinIF-Train for Model Adaptation

FinIF-Train provides the training component of FinIF. It contains 1,064 workflow-grounded examples across 38 task types and 7,851 annotated constraints. Unlike generic instruction-tuning data, each example is tied to a finance workflow, a concrete work product, and explicit delivery constraints. The intended use is supervised fine-tuning or preference-data construction for financial instruction following.

To avoid conflating dataset versions, we do not report legacy pilot SFT numbers in this draft. The same-split adaptation protocol is straightforward: train on FinIF-Train, evaluate both the base model and the adapted model on the 300-item FinIF-Test split, and report Micro IF, Macro item IF, Strict item pass rate, and auxiliary quality. This keeps model adaptation results directly comparable with the capability evaluation above.

6 Conclusion

We introduced FinIF, a benchmark suite and data construction framework for evaluating whether LLMs can produce usable financial work products under explicit delivery constraints. FinIF-Test provides a 300-item hard benchmark with 3,134 instruction-following constraints, while FinIF-Train provides 1,064 workflow-grounded examples for training and expansion. Experiments show that strong models satisfy most individual constraints but frequently fail strict item-level compliance. The largest weaknesses involve evidence placement, quantitative verification, and decision-boundary reasoning. These findings suggest that financial LLM evaluation should measure not only answer quality, but also whether the answer follows the operational instructions needed for auditability and downstream use.

Limitations

FinIF has several limitations. First, semantic evaluation depends on an LLM judge, so further human calibration and judge-sensitivity analysis are needed before treating scores as definitive. Second, some private financial records are synthetic, which protects privacy and improves controllability

but may not capture all institution-specific policies and data distributions. Third, the current reported experiments cover five completed model runs; broader evaluation over open-weight, finance-tuned, multilingual, and newly released models should be added. Fourth, FinIF-Test emphasizes hard instruction-following cases, so its strict pass rates should not be interpreted as general financial task success rates. Finally, same-split supervised fine-tuning experiments on the new 300-item test split remain future work for the next version of this draft.

Ethical Considerations

FinIF is designed for evaluating and improving instruction adherence in financial workflows, not for providing investment, legal, tax, or compliance advice. The benchmark should therefore be used as an offline evaluation and research resource rather than as an automated decision system. Synthetic private records are used to avoid exposing real customer or institutional data. For public-source contexts, users of the dataset should respect the licenses and terms of the underlying materials. Because higher instruction-following ability could make financial outputs appear more authoritative, deployed systems should still include human review, source verification, and institution-specific compliance controls.

References

Xin Guo, Haotian Xia, Zhaowei Liu, Hanyang Cao, Zhi Yang, Zhiqiang Liu, Sizhe Wang, Jinyi Niu, Chuqi Wang, Yanhui Wang, Xiaolong Liang, Xiaoming Huang, Bing Zhu, Zhongyu Wei, Yun Chen, Weining Shen, and Liwen Zhang. 2025. Fineval: A chinese financial domain knowledge evaluation benchmark for large language models. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 6258–6292. Association for Computational Linguistics.

Pranab Islam, Anand Kannappan, Douwe Kiela, Rebecca Qian, Nino Scherrer, and Bertie Vidgen. 2023. Financebench: A new benchmark for financial question answering. *arXiv preprint arXiv:2311.11944*.

Yuxin Jiang, Yufei Wang, Xingshan Zeng, Wanjun Zhong, Liangyou Li, Fei Mi, Lifeng Shang, Xin Jiang, Qun Liu, and Wei Wang. 2024. Followbench: A multi-level fine-grained constraints following benchmark for large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 4667–4688. Association for Computational Linguistics.

Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023. G-eval: Nlg evaluation using gpt-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2511–2522. Association for Computational Linguistics.

Guilong Lu, Xuntao Guo, Rongjunchen Zhang, Wenqiao Zhu, and Ji Liu. 2025. Bizfinbench: A business-driven real-world financial benchmark for evaluating llms. *arXiv preprint arXiv:2505.19457*.

Glenn Matlin, Siddharth, JM Anirudh, Aditya Shukla, Yahya Hassan, and Sudheer Chava. 2025. Financial instruction following evaluation (fife). *arXiv preprint arXiv:2512.08965*.

Ying Nie, Binwei Yan, Tianyu Guo, Hao Liu, Haoyu Wang, Wei He, Binfan Zheng, Weihao Wang, Qiang Li, Weijian Sun, Yunhe Wang, and Dacheng Tao. 2025. Cfinbench: A comprehensive chinese financial benchmark for large language models. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 876–891. Association for Computational Linguistics.

Yunjia Qi, Hao Peng, Xiaozhi Wang, Amy Xin, Youfeng Liu, Bin Xu, Lei Hou, and Juanzi Li. 2025. Agentif: Benchmarking instruction following of large language models in agentic scenarios. In *Advances in Neural Information Processing Systems*.

Yiwei Qin, Kaiqiang Song, Yebowen Hu, Wenlin Yao, Sangwoo Cho, Xiaoyang Wang, Xuansheng Wu, Fei Liu, Pengfei Liu, and Dong Yu. 2024. Infobench: Evaluating instruction following ability in large language models. In *Findings of the Association for Computational Linguistics: ACL 2024*. Association for Computational Linguistics.

Raj Sanjay Shah, Kunal Chawla, Dheeraj Eidnani, Agam Shah, Wendi Du, Sudheer Chava, Natraj Ramman, Charese Smiley, Jiaao Chen, and Diyi Yang. 2022. When flue meets flang: Benchmarks and large pretrained language model for financial domain. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2322–2335. Association for Computational Linguistics.

Encheng Su, Jianyu Wu, Chen Tang, Lintao Wang, Pengze Li, Aoran Wang, Jinouwen Zhang, Yizhou Wang, Yuan Meng, Xinzhu Ma, Shixiang Tang, and Houqiang Li. 2026. Sciif: Benchmarking scientific instruction following towards rigorous scientific intelligence. *arXiv preprint arXiv:2601.04770*.

Xiaodong Wu, Minhao Wang, Yichen Liu, Xiaoming Shi, He Yan, Lu Xiangju, Junmin Zhu, and Wei Zhang. 2025. Lifbench: Evaluating the instruction following performance and stability of large language models in long-context scenarios. In *Proceedings*

of the 63rd Annual Meeting of the Association for Computational Linguistics, pages 16445–16468. Association for Computational Linguistics.

Qianqian Xie, Weiguang Han, Zhengyu Chen, and 1 others. 2024. Finben: A holistic financial benchmark for large language models. In *Advances in Neural Information Processing Systems*.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena. *arXiv preprint arXiv:2306.05685*.

Jeffrey Zhou, Tianjian Lu, Swaroop Mishra, Siddhartha Brahma, Sujoy Basu, Yi Luan, Denny Zhou, and Le Hou. 2023. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*.

Appendix

A Dataset Details

FinIF-Train contains 1,064 workflow-grounded examples across 38 task types. Each task type contains 28 examples. FinIF-Test is selected as a 300-item hard benchmark from the same workflow design. The current FinIF-Test split contains 3,134 instruction-following constraints and 1,904 auxiliary diagnostic constraints. Auxiliary diagnostics are used for answer quality, finance-validity, and coverage analysis, but they are not included in the primary Micro IF, Macro IF, or Strict item pass metrics.

B Evaluation Details

The evaluator first applies deterministic rule checkers when a constraint is mechanically checkable. Rule checks cover length ranges, required labels, table schema, ordered-list counts, first-line formats, decimal precision, and related surface constraints. If a constraint cannot be decided reliably by local rules, it is sent to an item-batched LLM judge. The judge receives the source context, user query, model response, and a JSON schema listing all semantic constraints for the item. The judge must return a pass/fail decision and a short reason for each constraint. Outputs that do not parse as valid JSON are retried under the same schema. The reported experiments have 100% decision coverage on FinIF-Test.

C Training Data and Adaptation

FinIF-Train is intended for supervised fine-tuning and for constructing preference or verifier data. To keep future adaptation results comparable, models should be trained only on FinIF-Train and evaluated on the fixed 300-item FinIF-Test split. Legacy pilot results from earlier dataset versions are intentionally excluded from the main paper.

D Anonymized Release

For double-blind review, code and data should be released through an anonymized repository or anonymous sharing service. The camera-ready version can replace this with the public project URL after acceptance.